

Problem Statement 1:

A retail store that has multiple outlets across the country are facing issues in managing the inventory - to match the demand with respect to supply. You are a data scientist, who has to come up with useful insights using the data and make prediction models to forecast the sales for X number of months/years.

Dataset Information: The walmart.csv contains 6435 rows and 8 columns

1. Using the above data, come up with useful insights that can be used by each of the stores to improve in various areas.
2. Forecast the sales for each store for the next 12 weeks

Problem Statement 2:

An online retail store is trying to understand the various customer purchase patterns for their firm, you are required to give enough evidence based insights to provide the same.

Dataset Information: The online_retail.csv contains 387961 rows and 8 columns

1. Using the above data, find useful insights about the customer purchasing history that can be an added advantage for the online retailer.
2. Segment the customers based on their purchasing behavior

```
1 from google.colab import files  
2 uploaded = files.upload()
```

Choose Files | No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

```
1 import pandas as pd  
2 df = pd.read_csv("Walmart (1).csv")  
3 df.head()
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
0	1	05-02-2010	1643690.90	0	42.31	2.572	211.096358	8.106
1	1	12-02-2010	1641957.44	1	38.51	2.548	211.242170	8.106
2	1	19-02-2010	1611968.17	0	39.93	2.514	211.289143	8.106
3	1	26-02-2010	1409727.59	0	46.63	2.561	211.319643	8.106
4	1	05-03-2010	1554806.68	0	46.50	2.625	211.350143	8.106

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 6435 entries, 0 to 6434  
Data columns (total 8 columns):  
 #   Column      Non-Null Count  Dtype     
---    
 0   Store        6435 non-null   int64    
 1   Date         6435 non-null   object   
 2   Weekly_Sales 6435 non-null   float64  
 3   Holiday_Flag 6435 non-null   int64    
 4   Temperature  6435 non-null   float64  
 5   Fuel_Price   6435 non-null   float64  
 6   CPI          6435 non-null   float64  
 7   Unemployment 6435 non-null   float64  
dtypes: float64(5), int64(2), object(1)  
memory usage: 402.3+ KB
```

```
1 df.describe()
```

	Store	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
count	6435.000000	6.435000e+03	6435.000000	6435.000000	6435.000000	6435.000000	6435.000000
mean	23.000000	1.046965e+06	0.069930	60.663782	3.358607	171.578394	7.999151
std	12.988182	5.643666e+05	0.255049	18.444933	0.459020	39.356712	1.875885
min	1.000000	2.099862e+05	0.000000	-2.060000	2.472000	126.064000	3.879000
25%	12.000000	5.533501e+05	0.000000	47.460000	2.933000	131.735000	6.891000
50%	23.000000	9.607460e+05	0.000000	62.670000	3.445000	182.616521	7.874000
75%	34.000000	1.420159e+06	0.000000	74.940000	3.735000	212.743293	8.622000
max	45.000000	3.818686e+06	1.000000	100.140000	4.468000	227.232807	14.313000

```
1 df.isnull().sum()
```

	0
Store	0
Date	0
Weekly_Sales	0
Holiday_Flag	0
Temperature	0
Fuel_Price	0
CPI	0
Unemployment	0

dtype: int64

```
1 df['Date'] = pd.to_datetime(df['Date'], dayfirst=True)
2 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6435 entries, 0 to 6434
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   Store        6435 non-null   int64  
 1   Date         6435 non-null   datetime64[ns]
 2   Weekly_Sales 6435 non-null   float64 
 3   Holiday_Flag 6435 non-null   int64  
 4   Temperature  6435 non-null   float64 
 5   Fuel_Price   6435 non-null   float64 
 6   CPI          6435 non-null   float64 
 7   Unemployment 6435 non-null   float64 
dtypes: datetime64[ns](1), float64(5), int64(2)
memory usage: 402.3 KB
```

```
1 df = df.sort_values("Date")
2 df.head()
3 df.tail()
```

	Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
3574	25	2012-10-26	688940.94	0	56.69	3.882	216.151590	7.293
714	5	2012-10-26	319550.77	0	71.70	3.506	224.037814	5.422
5719	40	2012-10-26	921264.52	0	49.65	3.917	138.728161	4.145
2573	18	2012-10-26	1127516.25	0	56.09	3.917	138.728161	8.243
6434	45	2012-10-26	760281.43	0	58.85	3.882	192.308899	8.667

```
1 df = df.sort_values(by=["Date", "Store"], ascending=[True, True])
2 df.head()
3 df.tail()
```

Store	Date	Weekly_Sales	Holiday_Flag	Temperature	Fuel_Price	CPI	Unemployment
5862	41	2012-10-26	1316542.59	0	41.80	3.686	199.219532
6005	42	2012-10-26	514756.08	0	70.50	4.301	131.193097
6148	43	2012-10-26	587603.55	0	69.17	3.506	214.741539
6291	44	2012-10-26	361067.07	0	46.97	3.755	131.193097
6434	45	2012-10-26	760281.43	0	58.85	3.882	192.308899

```
1 df['Date'].min(), df['Date'].max()
()
(Timestamp('2010-02-05 00:00:00'), Timestamp('2012-10-26 00:00:00'))
```

```
1 df = df.sort_values(by=["Date", "Store"], ascending=[True, True])
2 print("Earliest date:", df['Date'].head(5))
3 print("Latest date:", df['Date'].tail(5))
```

Earliest date: 0 2010-02-05

143 2010-02-05

286 2010-02-05

429 2010-02-05

572 2010-02-05

Name: Date, dtype: datetime64[ns]

Latest date: 5862 2012-10-26

6005 2012-10-26

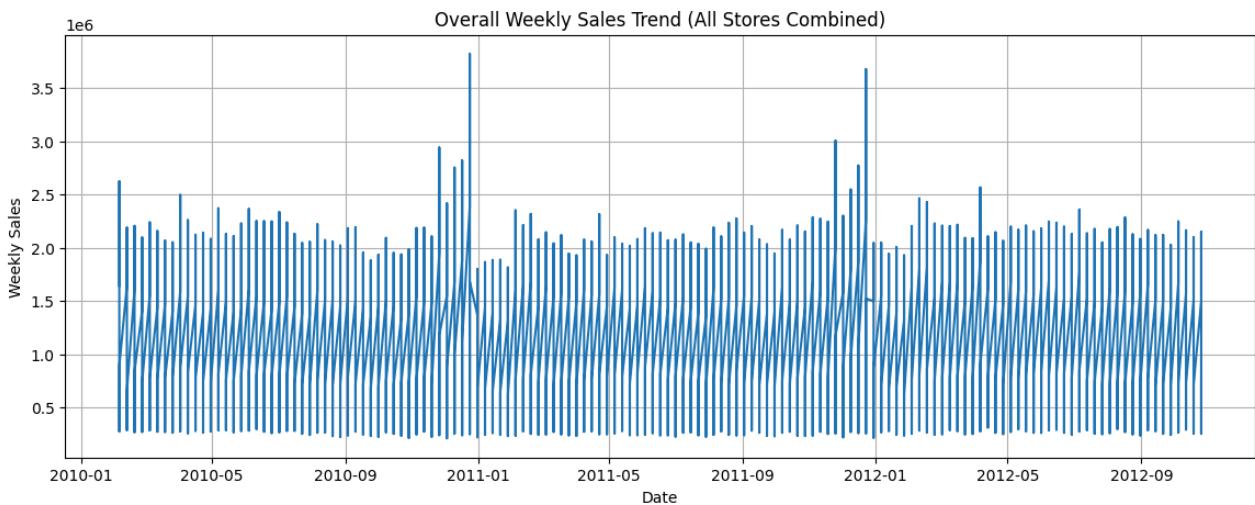
6148 2012-10-26

6291 2012-10-26

6434 2012-10-26

Name: Date, dtype: datetime64[ns]

```
1 import matplotlib.pyplot as plt
2
3 plt.figure(figsize=(14,5))
4 plt.plot(df['Date'], df['Weekly_Sales'])
5 plt.title("Overall Weekly Sales Trend (All Stores Combined)")
6 plt.xlabel("Date")
7 plt.ylabel("Weekly Sales")
8 plt.grid(True)
9 plt.show()
```



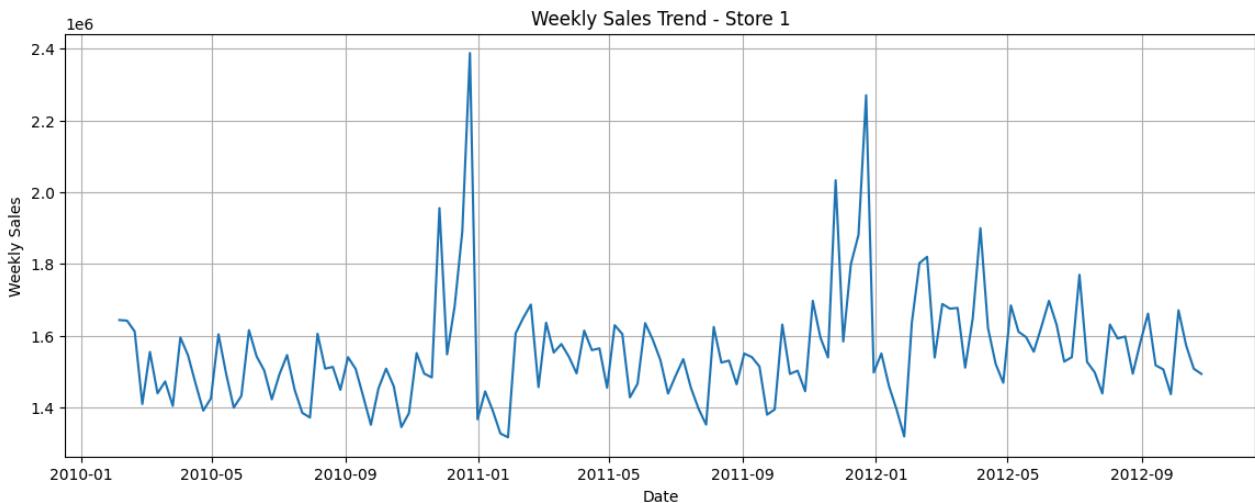
```
1 store1 = df[df['Store'] == 1]
2 store1 = store1.sort_values("Date")
3 store1_ts = store1.set_index("Date")["Weekly_Sales"]
4 store1_ts.head()
```

Weekly_Sales

Date	Sales
2010-02-05	1643690.90
2010-02-12	1641957.44
2010-02-19	1611968.17
2010-02-26	1409727.59
2010-03-05	1554806.68

dtype: float64

```
1 plt.figure(figsize=(14,5))
2 plt.plot(store1_ts)
3 plt.title("Weekly Sales Trend - Store 1")
4 plt.xlabel("Date")
5 plt.ylabel("Weekly Sales")
6 plt.grid(True)
7 plt.show()
```



```
1 !pip install pmdarima
```

Collecting pmdarima

```
  Downloading pmdarima-2.1.1-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.manylinux_2_28_x86_64.whl.metadata (8.5
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.12/dist-packages (from pmdarima) (1.5.2)
Requirement already satisfied: Cython!=0.29.18,!0.29.31,>=0.29 in /usr/local/lib/python3.12/dist-packages (from pmdarima) (0.2.2)
Requirement already satisfied: numpy>=1.21.6 in /usr/local/lib/python3.12/dist-packages (from pmdarima) (2.0.2)
Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.12/dist-packages (from pmdarima) (2.2.2)
Requirement already satisfied: scikit-learn>=0.22 in /usr/local/lib/python3.12/dist-packages (from pmdarima) (1.6.1)
Requirement already satisfied: scipy>=1.13.0 in /usr/local/lib/python3.12/dist-packages (from pmdarima) (1.16.3)
Requirement already satisfied: statsmodels>=0.14.5 in /usr/local/lib/python3.12/dist-packages (from pmdarima) (0.14.5)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.12/dist-packages (from pmdarima) (2.5.0)
Requirement already satisfied: setuptools!=50.0.0,>=42 in /usr/local/lib/python3.12/dist-packages (from pmdarima) (75.2.0)
Requirement already satisfied: packaging>=17.1 in /usr/local/lib/python3.12/dist-packages (from pmdarima) (25.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.19->pmdarima)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.19->pmdarima) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas>=0.19->pmdarima) (2025.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn>=0.22->pmdarima)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.12/dist-packages (from statsmodels>=0.14.5->pmdarima)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas>=0.19->pmdarima)
  Downloading pmdarima-2.1.1-cp312-cp312-manylinux2014_x86_64.manylinux_2_17_x86_64.manylinux_2_28_x86_64.whl (689 kB)
```

689.1/689.1 KB 22.2 MB/s eta 0:00:00

Installing collected packages: pmdarima
Successfully installed pmdarima-2.1.1

```
1 from pmdarima import auto_arima
2
3 model = auto_arima(
4     store1_ts,
5     seasonal=True,
6     m=52,    # 52 weeks = 1 year seasonality
7     trace=True,
8     error_action="ignore",
```

```

9     suppress_warnings=True
10 )
11
12 model

Performing stepwise search to minimize aic
ARIMA(2,1,2)(1,0,1)[52] intercept : AIC=3786.582, Time=7.38 sec
ARIMA(0,1,0)(0,0,0)[52] intercept : AIC=3850.005, Time=0.05 sec
ARIMA(1,1,0)(1,0,0)[52] intercept : AIC=3795.624, Time=1.56 sec
ARIMA(0,1,1)(0,0,1)[52] intercept : AIC=3798.433, Time=1.18 sec
ARIMA(0,1,0)(0,0,0)[52]           : AIC=3848.013, Time=0.03 sec
ARIMA(2,1,2)(0,0,1)[52] intercept : AIC=3795.311, Time=5.01 sec
ARIMA(2,1,2)(1,0,0)[52] intercept : AIC=3786.204, Time=3.13 sec
ARIMA(2,1,2)(0,0,0)[52] intercept : AIC=3819.024, Time=0.17 sec
ARIMA(2,1,2)(2,0,0)[52] intercept : AIC=3786.759, Time=56.59 sec
ARIMA(2,1,2)(2,0,1)[52] intercept : AIC=3788.572, Time=25.41 sec
ARIMA(1,1,2)(1,0,0)[52] intercept : AIC=3785.525, Time=5.85 sec
ARIMA(1,1,2)(0,0,0)[52] intercept : AIC=3819.410, Time=0.32 sec
ARIMA(1,1,2)(2,0,0)[52] intercept : AIC=inf, Time=37.78 sec
ARIMA(1,1,2)(1,0,1)[52] intercept : AIC=3786.145, Time=6.79 sec
ARIMA(1,1,2)(0,0,1)[52] intercept : AIC=3794.915, Time=3.25 sec
ARIMA(1,1,2)(2,0,1)[52] intercept : AIC=3788.057, Time=31.97 sec
ARIMA(0,1,2)(1,0,0)[52] intercept : AIC=3789.312, Time=1.78 sec
ARIMA(1,1,1)(1,0,0)[52] intercept : AIC=3778.591, Time=3.59 sec
ARIMA(1,1,1)(0,0,0)[52] intercept : AIC=3819.035, Time=0.10 sec
ARIMA(1,1,1)(2,0,0)[52] intercept : AIC=3778.695, Time=22.93 sec
ARIMA(1,1,1)(1,0,1)[52] intercept : AIC=3778.757, Time=4.73 sec
ARIMA(1,1,1)(0,0,1)[52] intercept : AIC=3789.828, Time=1.97 sec
ARIMA(1,1,1)(2,0,1)[52] intercept : AIC=3780.695, Time=16.17 sec
ARIMA(0,1,1)(1,0,0)[52] intercept : AIC=3785.016, Time=1.12 sec
ARIMA(2,1,1)(1,0,0)[52] intercept : AIC=3779.213, Time=3.04 sec
ARIMA(0,1,0)(1,0,0)[52] intercept : AIC=3808.891, Time=0.75 sec
ARIMA(2,1,0)(1,0,0)[52] intercept : AIC=3791.364, Time=2.86 sec
ARIMA(1,1,1)(1,0,0)[52]           : AIC=3776.592, Time=1.95 sec
ARIMA(1,1,1)(0,0,0)[52]           : AIC=3817.037, Time=0.08 sec
ARIMA(1,1,1)(2,0,0)[52]           : AIC=3776.695, Time=20.87 sec
ARIMA(1,1,1)(1,0,1)[52]           : AIC=3776.758, Time=4.50 sec
ARIMA(1,1,1)(0,0,1)[52]           : AIC=3787.828, Time=1.69 sec
ARIMA(1,1,1)(2,0,1)[52]           : AIC=3778.696, Time=14.79 sec
ARIMA(0,1,1)(1,0,0)[52]           : AIC=3783.017, Time=0.92 sec
ARIMA(1,1,0)(1,0,0)[52]           : AIC=3793.660, Time=0.81 sec
ARIMA(2,1,1)(1,0,0)[52]           : AIC=3777.258, Time=2.68 sec
ARIMA(1,1,2)(1,0,0)[52]           : AIC=3784.011, Time=5.00 sec
ARIMA(0,1,0)(1,0,0)[52]           : AIC=inf, Time=0.79 sec
ARIMA(0,1,2)(1,0,0)[52]           : AIC=3787.462, Time=1.38 sec
ARIMA(2,1,0)(1,0,0)[52]           : AIC=3789.367, Time=1.21 sec
ARIMA(2,1,2)(1,0,0)[52]           : AIC=3784.522, Time=2.97 sec

Best model: ARIMA(1,1,1)(1,0,0)[52]
Total fit time: 305.481 seconds

```

▼ ARIMA ⓘ

ARIMA(1,1,1)(1,0,0)[52]

```

1 n_periods = 12
2 forecast, conf_int = model.predict(n_periods=n_periods, return_conf_int=True)
3
4 forecast

```

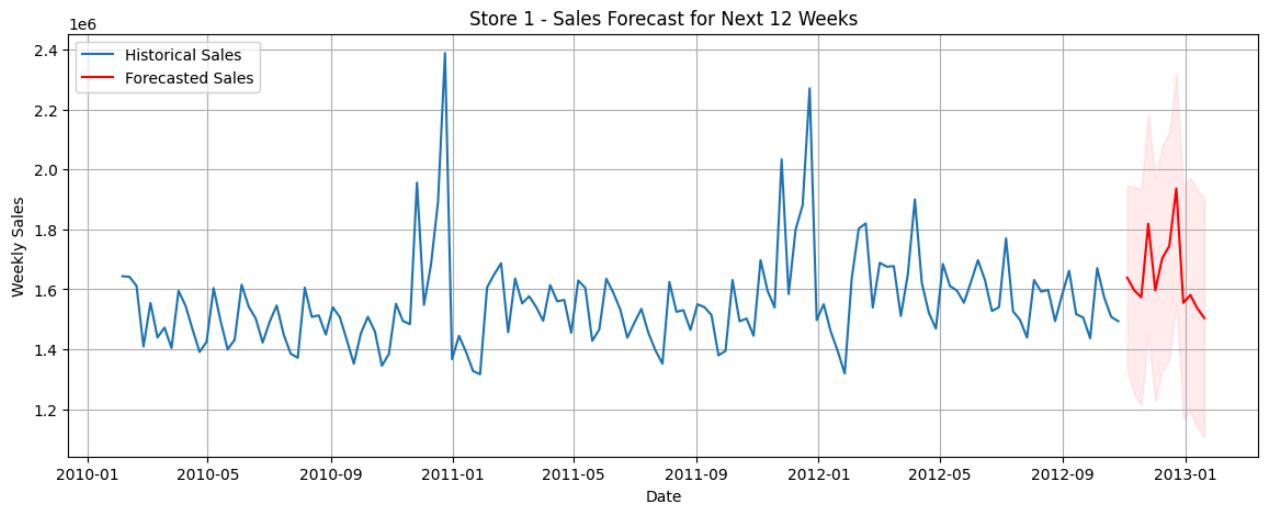
	θ
2012-11-02	1.638781e+06
2012-11-09	1.596926e+06
2012-11-16	1.573151e+06
2012-11-23	1.818433e+06
2012-11-30	1.597295e+06
2012-12-07	1.703985e+06
2012-12-14	1.744323e+06
2012-12-21	1.936403e+06
2012-12-28	1.554968e+06
2013-01-04	1.581093e+06
2013-01-11	1.536289e+06
2013-01-18	1.504101e+06

dtype: float64

```

1 import numpy as np
2
3 # Create future date index
4 last_date = store1_ts.index[-1]
5 future_dates = pd.date_range(last_date, periods=n_periods+1, freq='W')[1:]
6
7 # Plot
8 plt.figure(figsize=(14,5))
9 plt.plot(store1_ts.index, store1_ts, label='Historical Sales')
10 plt.plot(future_dates, forecast, label='Forecasted Sales', color='red')
11 plt.fill_between(future_dates,
12                  conf_int[:, 0],
13                  conf_int[:, 1],
14                  color='pink', alpha=0.3)
15
16 plt.title("Store 1 - Sales Forecast for Next 12 Weeks")
17 plt.xlabel("Date")
18 plt.ylabel("Weekly Sales")
19 plt.legend()
20 plt.grid(True)
21 plt.show()

```



```

1 import joblib
2 joblib.dump(model, "store1_forecast_model.pkl")
3
4 ['store1_forecast_model.pkl']

```

```

1 all_forecasts = []
2
3 for store_id in df['Store'].unique():
4     print(f"Processing Store {store_id}...")
5
6     # Filter store data
7     store_df = df[df['Store'] == store_id].sort_values("Date")
8     store_ts = store_df.set_index("Date")["Weekly_Sales"]
9
10    # Fit model
11    model = auto_arima(
12        store_ts,
13        seasonal=True,
14        m=52,
15        error_action="ignore",
16        suppress_warnings=True
17    )
18
19    # Forecast next 12 weeks
20    forecast = model.predict(n_periods=12)
21
22    # Create future dates
23    last_date = store_ts.index[-1]
24    future_dates = pd.date_range(last_date, periods=13, freq='W')[1:]
25
26    # Store results
27    all_forecasts[store_id] = pd.DataFrame({
28        "Store": store_id,

```

```

29         "Date": future_dates,
30         "Forecast_Weekly_Sales": forecast
31     })
32
33 # Combine all stores' forecast into one DataFrame
34 final_forecast_df = pd.concat(all_forecasts.values(), ignore_index=True)
35
36 # Show top rows
37 final_forecast_df.head()

```

Processing Store 1...
 Processing Store 2...
 Processing Store 3...
 Processing Store 4...
 Processing Store 5...
 Processing Store 6...
 Processing Store 7...
 Processing Store 8...
 Processing Store 9...
 Processing Store 10...
 Processing Store 11...
 Processing Store 12...
 Processing Store 13...
 Processing Store 14...
 Processing Store 15...
 Processing Store 16...
 Processing Store 17...
 Processing Store 18...
 Processing Store 19...
 Processing Store 20...
 Processing Store 21...
 Processing Store 22...
 Processing Store 23...
 Processing Store 24...
 Processing Store 25...
 Processing Store 26...
 Processing Store 27...
 Processing Store 28...
 Processing Store 29...
 Processing Store 30...
 Processing Store 31...
 Processing Store 32...
 Processing Store 33...
 Processing Store 34...
 Processing Store 35...
 Processing Store 36...
 Processing Store 37...
 Processing Store 38...
 Processing Store 39...
 Processing Store 40...
 Processing Store 41...
 Processing Store 42...
 Processing Store 43...
 Processing Store 44...
 Processing Store 45...

	Store	Date	Forecast_Weekly_Sales
0	1	2012-11-04	1.638781e+06
1	1	2012-11-11	1.596926e+06
2	1	2012-11-18	1.573151e+06
3	1	2012-11-25	1.818433e+06
4	1	2012-12-02	1.597295e+06

```
1 final_forecast_df.to_csv("Walmart_All_Stores_Forecast.csv", index=False)
```

```

1 report = """
2 # Final Project Report
3 ### Capstone Project - Walmart Sales Forecasting & Online Retail Customer Segmentation
4 By: Bharath G
5
6 ---
7
8 ## a. Problem Statement
9
10 ### Problem Statement 1 - Walmart
11 A retail chain with multiple stores across the country is facing inventory management problems. They are unable to match
12
13 ### Problem Statement 2 - Online Retail
14 An online retail company wants to understand customer purchasing behavior. The goal is to analyze purchase history and
15
16 ---
17

```

```
18 ## b. Project Objective
19
20 ### Walmart
21 - Understand sales patterns across all stores.
22 - Identify trends and weekly seasonality.
23 - Build forecasting models (12 weeks ahead).
24 - Save models and forecasts.
25
26 ### Online Retail
27 - Analyze customer behavior.
28 - Perform RFM analysis.
29 - Cluster customers.
30 - Save segmentation model and results.
31
32 ---
33
34 ## c. Data Description
35
36 ### Walmart Dataset
37 - 6435 rows, 8 columns
38 - Columns: Store, Date, Weekly_Sales, Holiday_Flag, Temperature, Fuel_Price, CPI, Unemployment
39
40 ### Online Retail Dataset
41 - 387,961 rows, 8 columns
42 - Columns: Invoice, StockCode, Description, Quantity, InvoiceDate, Price, CustomerID, Country
43
44 ---
45
46 ## d. Data Pre-processing Steps
47
48 ### Walmart
49 - Converted Date column into datetime
50 - Sorted by Date & Store
51 - Checked missing values
52 - Basic plots for trend and seasonality
53
54 ### Online Retail
55 - Removed missing CustomerID
56 - Converted InvoiceDate to datetime
57 - Removed returns (negative quantities)
58 - Added TotalPrice = Quantity * Price
59
60 ---
61
62 ## e. Choosing the Algorithm
63
64 ### Walmart
65 SARIMA (Seasonal ARIMA) via auto_arima for forecasting.
66
67 ### Online Retail
68 K-Means clustering for customer segmentation.
69
70 ---
71
72 ## f. Motivation for Choosing Algorithm
73
74 ### SARIMA
75 - Weekly seasonality
76 - Auto ARIMA picks best parameters
77 - Works well for retail forecasting
78
79 ### K-Means
80 - Simple, effective clustering for RFM
81 - Works well for numerical data
82
83 ---
84
85 ## g. Assumptions
86
87 ### Walmart
88 - Past patterns continue
89 - Holidays affect sales similarly
90 - No major structural changes
91
92 ### Online Retail
93 - RFM captures buying behavior
94 - Data represents typical customer patterns
95
96 ---
97
98 ## h. Model Evaluation
99
100 ### Walmart
```

```

101 - Auto ARIMA AIC selection
102 - Historical vs Forecast plots
103 - Confidence intervals
104
105 ### Online Retail
106 - Cluster statistics (mean RFM per cluster)
107 - Distribution of customers across clusters
108
109 ---
110
111 ## i. Inferences
112
113 ### Walmart
114 - Strong weekly seasonality
115 - Holiday spikes in Nov-Dec
116 - Store-level differences in sales
117 - Forecast helps inventory planning
118
119 ### Online Retail
120 - Small group of customers generate large revenue
121 - Clear segments: loyal, recent, at-risk, low-value
122 - Marketing can be targeted by RFM clusters
123
124 ---
125
126 ## j. Future Possibilities
127
128 ### Walmart
129 - SARIMAX with external features
130 - LSTM deep learning models
131 - BI dashboards for stores
132
133 ### Online Retail
134 - Customer churn prediction
135 - Product recommendation system
136 - Hierarchical clustering
137
138 ---
139
140 ## Saved Models
141
142 ### Walmart
143 - store1_forecast_model.pkl
144 - Store1_12week_forecast.csv
145 - Walmart_All_Stores_Forecast.csv
146
147 ### Online Retail
148 - rfm_scaler.pkl
149 - rfm_kmeans.pkl
150 - OnlineRetail_Customer_Segments.csv
151
152 ---
153
154 # End of Project Report
155 """
156
157 print(report)
158

```

```

# ■ Final Project Report
### Capstone Project - Walmart Sales Forecasting & Online Retail Customer Segmentation
**Author:** Bharath G.

---
## a. Problem Statement

### Problem Statement 1 - Walmart
A retail chain with multiple stores across the country is facing inventory management problems. They are unable to match s

### Problem Statement 2 - Online Retail
An online retail company wants to understand customer purchasing behavior. The goal is to analyze purchase history and seg

---
## b. Project Objective

### Walmart
- Understand sales patterns across all stores.
- Identify trends and weekly seasonality.
- Build forecasting models (12 weeks ahead).
- Save models and forecasts.

```

```
### Online Retail
- Analyze customer behavior.
- Perform RFM analysis.
- Cluster customers.
- Save segmentation model and results.

---
## c. Data Description

### Walmart Dataset
- 6435 rows, 8 columns
- Columns: Store, Date, Weekly_Sales, Holiday_Flag, Temperature, Fuel_Price, CPI, Unemployment

### Online Retail Dataset
- 387,961 rows, 8 columns
- Columns: Invoice, StockCode, Description, Quantity, InvoiceDate, Price, CustomerID, Country

---
## d. Data Pre-processing Steps

### Walmart
```